



ADVERSARY
ARGUMENT

Perhaps you find the term “accidentally” unconvincing—after all, there might be an algorithm that happens not to take the dead-end path as it explores. Our claim, to be more precise, is that *no algorithm can avoid dead ends in all state spaces*. Consider the two dead-end state spaces in Figure 4.19(a). To an online search algorithm that has visited states *S* and *A*, the two state spaces look *identical*, so it must make the same decision in both. Therefore, it will fail in one of them. This is an example of an **adversary argument**—we can imagine an adversary that constructs the state space while the agent explores it and can put the goals and dead ends wherever it likes.