


**Figure 4.18** A simple maze problem. The agent starts at  $S$  and must reach  $G$ , but knows nothing of the environment.



**Figure 4.19** (a) Two state spaces that might lead an online search agent into a dead end. Any given agent will fail in at least one of these spaces. (b) A two-dimensional environment that can cause an online search agent to follow an arbitrarily inefficient route to the goal. Whichever choice the agent makes, the adversary blocks that route with another long, thin wall, so that the path followed is much longer than the best possible path.



ADVERSARY  
ARGUMENT

Perhaps you find the term “accidentally” unconvincing—after all, there might be an algorithm that happens not to take the dead-end path as it explores. Our claim, to be more precise, is that *no algorithm can avoid dead ends in all state spaces*. Consider the two dead-end state spaces in Figure 4.19(a). To an online search algorithm that has visited states  $S$  and  $A$ , the two state spaces look *identical*, so it must make the same decision in both. Therefore, it will fail in one of them. This is an example of an **adversary argument**—we can imagine an adversary that constructs the state space while the agent explores it and can put the goals and dead ends wherever it likes.