


Solving problems by searching

CS4881 - Artificial Intelligence
Jay Urbain, Ph.D.



Outline

- Problem-solving agents
- Problem types
- Problem formulation
- Example problems
- Basic search algorithms



Reflex & Problem Solving Agents

- Reflex agents

- Base their actions on a direct mapping from states to actions.
- Cannot operate well in environments requiring a sequence of operations
 - mapping sequences of state-actions would be too large to store, and would take too long to learn.

- Problem Solving or Goal based agents

- Can succeed in this environment by considering future actions and the desirability of their outcomes.
- A Problem solving agent is a goal based-agent that use atomic representations.
 - where states of the world are considered as wholes (no internal structure is visible to the agent).



Problem Solving Agents

- How an agent can find a *sequence of actions that achieves its goals*, when no single action will do.
- Goals help organize behavior by *limiting the courses of action* an agent seeks to achieve.
- *Goal formulation* based on the *current state* and the agent's performance measure is the 1st step in problem solving.
- *Problem formulation* is the process of deciding what actions and states to consider, given a goal.



Travel from Arad to Bucharest

- Map of Romania
 - Agent with several options of unknown value can decide what to do by first examining future actions that eventually lead to states of known value.
 - *Environment?*



Travel from Arad to Bucharest

- Environment assumptions:
 - Observable (know current state – *have map*)
 - Discrete (finite number of actions to choose)
 - Known (which states are reached by each action).
 - Deterministic (each action has one outcome).
- Under these assumptions:
 - *Solution to any problem is a fixed sequence of actions.*



Search

- *Search: sequence of actions to achieve goal.*
 - Search algorithm takes a problem as input and returns a solution in the form of a sequence of actions.
 - Once solution is found, actions it recommends are executed.

Formulate -> Search -> Action

- *When executing, the agent is running open loop, i.e., it ignores percepts since it already knows in advance what they will be.*



Formulate Problem

- State space of problem – forms directed graph
 - Initial state
 - Description of the possible actions
 - Description of what each action does (transition model)
- Goal test
- Path cost



Problem-solving agents

```
function SIMPLE-PROBLEM-SOLVING-AGENT(percept) returns an action
  static: seq, an action sequence, initially empty
           state, some description of the current world state
           goal, a goal, initially null
           problem, a problem formulation

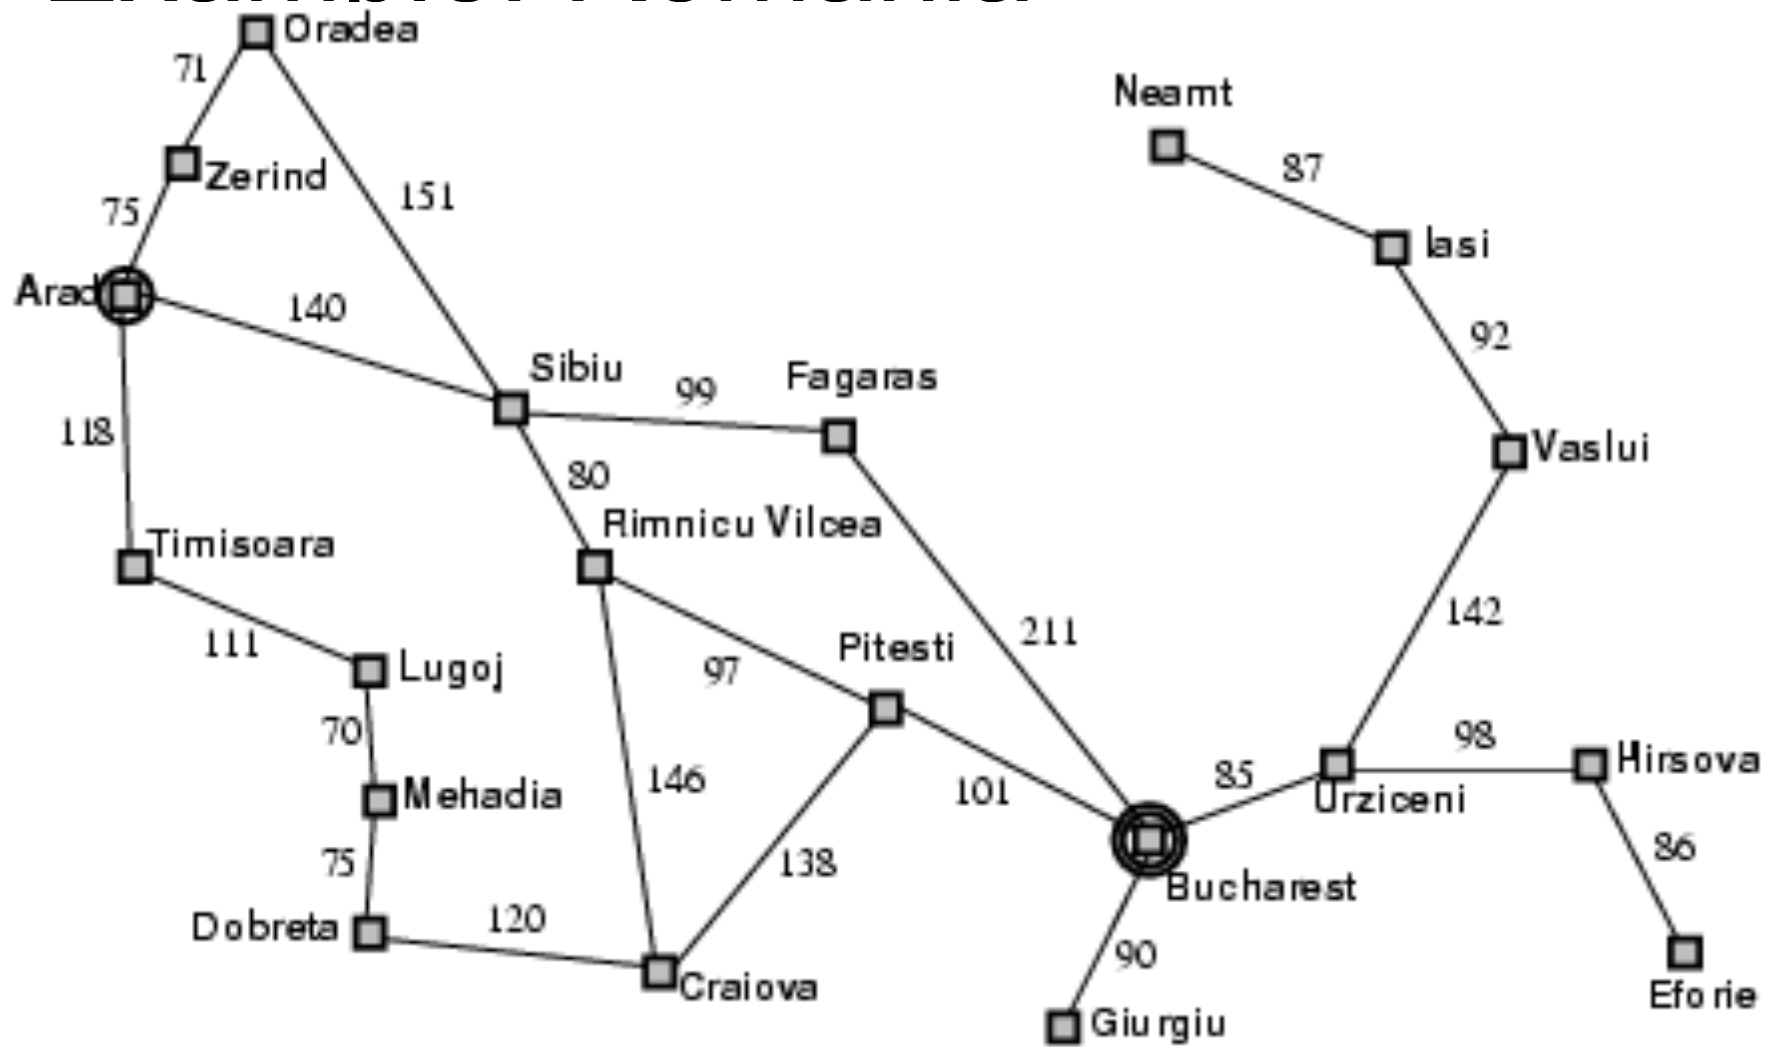
  state  $\leftarrow$  UPDATE-STATE(state, percept)
  if seq is empty then do
    goal  $\leftarrow$  FORMULATE-GOAL(state)
    problem  $\leftarrow$  FORMULATE-PROBLEM(state, goal)
    seq  $\leftarrow$  SEARCH(problem)
  action  $\leftarrow$  FIRST(seq)
  seq  $\leftarrow$  REST(seq)
  return action
```



Example: Romania

- On spring break in Romania; currently in Arad.
- Flight leaves tomorrow from Bucharest to Milwaukee.
- Formulate goal:
 - be in Bucharest
- Formulate problem (what actions and states to consider, given initial state and a goal):
 - **states**: various cities
 - **actions**: drive between cities
- Find solution:
 - sequence of cities, e.g., Arad, Sibiu, Fagaras, Bucharest

Example: Romania






Problem types

- Deterministic, fully observable → single-state problem
 - Agent knows exactly which state it will be in; *solution is a sequence.*
- Non-observable → sensorless problem (conformant problem)
 - Agent may have no idea where it is; *solution is a sequence.*
- Nondeterministic and/or partially observable → contingency problem
 - percepts provide new information about current state
 - often interleave search and, execution
- Unknown state space → exploration problem


Example: vacuum world

- Single-state (fully observable, deterministic), start in #5. Solution?




Example: vacuum world

- Single-state, start in #5.
Solution? [*Right, Suck*]
- Sensorless (non-observable), start in {1,2,3,4,5,6,7,8} e.g.,
Right goes to {2,4,6,8}
Solution?



Example: vacuum world

- **Sensorless**, start in {1,2,3,4,5,6,7,8} e.g.,
Right goes to {2,4,6,8}
Solution? [*Right*, *Suck*, *Left*, *Suck*]
- **Contingency**
(Nondeterministic and/or partially observable)
 - Nondeterministic: *Suck* may dirty a clean carpet
 - Partially observable: location, dirt at current location.
 - Percept: [*L*, *Clean*], i.e., start in #5 or #7
Solution?




Example: vacuum world

■ Contingency

- Nondeterministic: *Suck* may dirty a clean carpet
- Partially observable: location, dirt at current location.
- Percept: $[L, \text{Clean}]$, i.e., start in #5 or #7

Solution? $[Right, \text{if dirt then Suck}]$





State problem formulation

A **problem** is defined by four items:


1. **initial state** e.g., "at Arad"
 2. **actions** or **successor function** $S(x)$ = set of action–state pairs
 - e.g., $S(\text{Arad}) = \{ \langle \text{Arad} \rightarrow \text{Zerind}, \text{Sibiu} \rangle, \dots \}$
 3. **goal test**, can be
 - **explicit**, e.g., $x = \text{"at Bucharest"}$
 - **implicit**, e.g., $\text{Checkmate}(x)$
 4. **path cost** (additive)
 - e.g., sum of distances, number of actions executed, etc.
 - $c(x, a, y)$ is the **step cost**, assumed to be ≥ 0
- A **solution** is a sequence of actions leading from the initial state to a goal state



Selecting a state space


- Real world is absurdly complex
 - state space must be **abstracted** for problem solving
- (Abstract) state = set of real states
- (Abstract) action = complex combination of real actions
 - e.g., "Arad → Zerind" represents a complex set of possible routes, detours, rest stops, etc.
- **(Abstract) solution =**
 - ***set of real paths that are solutions in the real world***
 - Each abstract action should be "easier" than the original problem

Vacuum world state space graph



- states?
- actions?
- goal test?
- path cost?

Vacuum world state space graph



- states? dirt and robot location
- actions? *Left, Right, Suck*
- goal test? no dirt at all locations
- path cost? 1 per action

Example: The 8-puzzle

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- states?
- actions?
- goal test?
- path cost?

Example: The 8-puzzle

7	2	4
5		6
8	3	1

Start State


	1	2
3	4	5
6	7	8

Goal State

- states? locations of tiles ($9! = 362,880$)
- actions? move blank left, right, up, down
- goal test? = goal state (given)
- path cost? 1 per move

[Note: optimal solution of n -Puzzle family is NP-hard]

Example: robotic assembly



- states?: real-valued coordinates of robot joint angles
parts of the object to be assembled
- actions?: continuous motions of robot joints
- goal test?: complete assembly
- path cost?: time to execute




Tree search algorithms

- Basic idea:
 - offline, simulated exploration of state space by generating successors of already-explored states (a.k.a. ~expanding states)


```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
```


Tree search example






Tree search example



Tree search example




Implementation: general tree search

```
function TREE-SEARCH(problem, fringe) returns a solution, or failure  
  fringe  $\leftarrow$  INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)  
  loop do  
    if fringe is empty then return failure  
    node  $\leftarrow$  REMOVE-FRONT(fringe)  
    if GOAL-TEST[problem](STATE[node]) then return SOLUTION(node)  
    fringe  $\leftarrow$  INSERTALL(EXPAND(node, problem), fringe)
```

```
function EXPAND(node, problem) returns a set of nodes  
  successors  $\leftarrow$  the empty set  
  for each action, result in SUCCESSOR-FN[problem](STATE[node]) do  
    s  $\leftarrow$  a new NODE  
    PARENT-NODE[s]  $\leftarrow$  node; ACTION[s]  $\leftarrow$  action; STATE[s]  $\leftarrow$  result  
    PATH-COST[s]  $\leftarrow$  PATH-COST[node] + STEP-COST(node, action, s)  
    DEPTH[s]  $\leftarrow$  DEPTH[node] + 1  
    add s to successors  
  return successors
```

Implementation: states vs. nodes

- A **state** is a (representation of) a physical configuration
- A **node** is a data structure constituting part of a search tree includes **state**, **parent node**, **action**, **path cost** $g(x)$, **depth**



- The `Expand` function creates new nodes, filling in the various fields and using the `SuccessorFn` of the problem to create the corresponding states.



Search strategies

- A search strategy is defined by picking the **order of node expansion**
- Strategies are evaluated along the following dimensions:
 - **completeness**: does it always find a solution if one exists?
 - **time complexity**: number of nodes generated
 - **space complexity**: maximum number of nodes in memory
 - **optimality**: does it always find a least-cost solution?
- Time and space complexity are measured in terms of
 - *b*: maximum branching factor of the search tree
 - *d*: depth of the least-cost solution
 - *m*: maximum depth of the state space (may be ∞)




Uninformed search strategies

Uninformed search strategies use only the information available in the problem definition

- Breadth-first search
- Uniform-cost search
- Depth-first search
- Depth-limited search
- Iterative deepening search


Breadth-first search

- Expand shallowest unexpanded node
- Implementation:
 - *fringe* is a FIFO queue, i.e., new successors go at end




Breadth-first search

- Expand shallowest unexpanded node
- Implementation:
 - *fringe* is a FIFO queue, i.e., new successors go at end




Breadth-first search

- Expand shallowest unexpanded node
- Implementation:
 - *fringe* is a FIFO queue, i.e., new successors go at end



Breadth-first search

- Expand shallowest unexpanded node
- Implementation:
 - *fringe* is a FIFO queue, i.e., new successors go at end





Breadth-first search

```
enqueue the root node // FIFO Queue
while(true)
    dequeue a node and examine it.
    if goal(node)
        return result
    else
        enqueue successor nodes (direct child nodes).
    if queue is empty
        return "not found"
```



Properties of breadth-first search

- Complete?
- Time?
- Space?
- Optimal?



Properties of breadth-first search

- Complete? Yes (if b is finite)
- Time? $1+b+b^2+b^3+\dots +b^d = O(b^d)^*$
- Space? $O(b^d)^*$ (keeps every node in memory)
- Optimal? Yes (if cost = 1 per step)

- **Space** is the bigger problem (more than time)!

**Note: $O(b^{d+1})$ if goal test applied to nodes when selected for expansion rather than when generated, the whole layer of nodes at depth d would be expanded before the goal was detected.*




Uniform-cost search

- Expand least-cost unexpanded node
- Implementation:
 - *fringe* = queue ordered by path cost
- Equivalent to breadth-first if step costs all equal
- Complete? Yes, if step cost $\geq \epsilon$
- Time? # of nodes with $g \leq$ cost of optimal solution,
 $O(b^{\text{ceiling}(C^*/\epsilon)})$ where C^* is the cost of the optimal solution
- Space? # of nodes with $g \leq$ cost of optimal solution,
 $O(b^{\text{ceiling}(C^*/\epsilon)})$
- Optimal? Yes – nodes expanded in increasing order of $g(n)$


Depth-first search

- Expand deepest unexpanded node
- Implementation:
 - *fringe* = LIFO queue (stack), i.e., put successors at front



Depth-first search

- Expand deepest unexpanded node
- Implementation:
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- Implementation:
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- Implementation:
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- Implementation:
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- Implementation:
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- Implementation:
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- Implementation:
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- Implementation:
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- Implementation:
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- Implementation:
 - *fringe* = LIFO queue, i.e., put successors at front



Depth-first search

- Expand deepest unexpanded node
- Implementation:
 - *fringe* = LIFO queue, i.e., put successors at front





Properties of depth-first search

- Complete?
- Time?
- Space?
- Optimal?



Properties of depth-first search

- Complete? No: fails in infinite-depth spaces, spaces with loops
 - Modify to avoid repeated states along path
 - complete in finite spaces
- Time? $O(b^m)$: terrible if m is much larger than d
 - but if solutions are dense, may be much faster than breadth-first
- Space? $O(bm)$, i.e., linear space!
- Optimal? No



Depth-first search (iterative)

```
push the root node // LIFO Queue == stack
while(true)
    pop a node and examine it.
    if goal(node)
        return result
    else
        push successor nodes (direct child nodes).
if queue is empty
    return "not found"
```



Depth-first search (recursive)

```
dfs(node n)
  if goal(n)
    return n
  for each child node c
    dfs(c);
```



Depth-limited search

= depth-first search with depth limit l ,
i.e., nodes at depth l have no successors

■

```
function DEPTH-LIMITED-SEARCH(problem, limit) returns soln/fail/cutoff
  RECURSIVE-DLS(MAKE-NODE(INITIAL-STATE[problem]), problem, limit)

function RECURSIVE-DLS(node, problem, limit) returns soln/fail/cutoff
  cutoff-occurred? ← false
  if GOAL-TEST[problem](STATE[node]) then return SOLUTION(node)
  else if DEPTH[node] = limit then return cutoff
  else for each successor in EXPAND(node, problem) do
    result ← RECURSIVE-DLS(successor, problem, limit)
    if result = cutoff then cutoff-occurred? ← true
    else if result ≠ failure then return result
  if cutoff-occurred? then return cutoff else return failure
```




Iterative deepening search

```
function ITERATIVE-DEEPENING-SEARCH( problem) returns a solution, or fail-  
ure  
  inputs: problem, a problem  
  for depth  $\leftarrow$  0 to  $\infty$  do  
    result  $\leftarrow$  DEPTH-LIMITED-SEARCH( problem, depth)  
    if result  $\neq$  cutoff then return result
```



Iterative deepening search / =0

Limit = 0




Iterative deepening search / =1

Limit = 1




Iterative deepening search / =2

Limit = 2



Iterative deepening search / =3

Limit = 3





Iterative deepening search

- Number of nodes generated in an iterative deepening search to depth d with branching factor b :

$$N_{IDS} = d \cdot b^1 + (d-1)b^2 + \dots + 3b^{d-2} + 2b^{d-1} + 1b^d$$

- For $b = 10$, $d = 5$,
 - $N_{BFS} = 10 + 100 + 1,000 + 10,000 + 100,000 = 111,111$
 - $N_{IDS} = 50 + 400 + 3,000 + 20,000 + 100,000 = 123,456$
- Overhead = $(123,456 - 111,111)/111,111 = 11\%$



Properties of iterative deepening search

- Complete? Yes
- Time? $(d+1)b^0 + d b^1 + (d-1)b^2 + \dots + b^d = O(b^d)$
- Space? $O(bd)$
- Optimal? Yes, if step cost = 1





Summary of algorithms

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening
Complete?	Yes	Yes	No	No	Yes
Time	$O(b^{d+1})$	$O(b^{\lceil C^*/\epsilon \rceil})$	$O(b^m)$	$O(b^l)$	$O(b^d)$
Space	$O(b^{d+1})$	$O(b^{\lceil C^*/\epsilon \rceil})$	$O(bm)$	$O(bl)$	$O(bd)$
Optimal?	Yes	Yes	No	No	Yes

Repeated states


- Failure to detect repeated states can turn a linear problem into an exponential one!





Graph search

```
function GRAPH-SEARCH(problem, fringe) returns a solution, or failure
  closed ← an empty set
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST[problem](STATE[node]) then return SOLUTION(node)
    if STATE[node] is not in closed then
      add STATE[node] to closed
      fringe ← INSERTALL(EXPAND(node, problem), fringe)
```



Summary (1)

- Actions in environments that are *deterministic*, *observable*, *static*, and *completely known*.
- In these cases, an agent can construct sequences of actions that achieve its goals – this process is called *search*.
- Problem formulation usually requires *abstracting* away real-world details to define a state space that can feasibly be explored.
- Variety of uninformed search strategies.
- Iterative deepening search uses only linear space and not much more time than other uninformed algorithms.



Summary (2)

- Before an agent can start searching for solutions, it must formulate a **goal**, then use the goal to formulate a **problem**.
- A problem consists of four parts:
 1. Initial state
 2. Set of states
 3. Set of actions
 4. Goal test function
 5. Path cost function
- A **path** through the **state space** from the **initial state** to a **goal state** is a **solution**.
- A single general TREE-SEARCH algorithm can be used to solve any problem; specific variants of the algorithm embody different strategies.



Summary (3)

- Search algorithms are judged on the basis of:
 - Completeness
 - Optimality
 - Time complexity
 - Space complexity
- Breadth-first search – complete, optimal for unit step costs, and has space & time complexity costs of $O(b^d)$. *Where d is depth of shallowest soln.*
- Uniform-cost search – like Breadth-first search, but expands node with lowest cost. It's complete and optimal if cost of each step exceeds some positive ϵ .
- Depth-first search – selects the deepest unexpanded node in the search tree for expansion. It's neither complete nor optimal has time complexity costs of $O(b^m)$ and space complexity costs $O(bd)$. *Where m is max depth.*
- Depth-limited search imposes a fixed limit on depth-first search.



Summary (4)

- Bi-directional search can enormously reduce time complexity, but is not always applicable and may require too much space.
- For graphs it is usually important to check for repeated states. GRAPH-SEARCH algorithm checks for repeated states.
- When the environment is partially observable, the agent can apply search algorithms in the space of belief states, or sets of states the agent might be in.
- A contingency plan is needed to handle unknown circumstances that may arise.