


# Adversarial Search

## Minimax

CS4881 Artificial Intelligence

Jay Urbain, Ph.D.

# Game Theory



# Dilbert Teaches Game Theory



# Outline

- Optimal decisions
- Minimax
- $\alpha$ - $\beta$  pruning
- Imperfect, real-time decisions

# Adversarial Search

- Multiagent environments
  - Any given agent needs to consider the actions of other agents and how they affect their welfare.
- Unpredictability of other agents can introduce many possible contingencies into the agent's problem-solving process.
- Distinguish *cooperative* and *competitive* multi-agent environments.
- Competitive environments in which the agents' goals are in conflict - give rise to *adversarial search problems* – often called *games*.

# Games in AI

- Game theory (Economics) – views multiagent environments a game regardless of whether the agents are *cooperative* or *competitive*.
- In AI, “games” are usually:
  - Deterministic
  - Turn-taking
  - Two-player
  - Zero-sum games of perfection
- In our agent environment framework this means:
  - Deterministic (Strategic)
  - Fully observable
  - Multiagent (2) where actions must alternate
  - Utility values at end of game are equal and opposite {win,loose}, {1,0}

# Games vs. search problems

- **Unpredictable** opponent → specifying a move for every possible opponent reply.
- **Time limits** → unlikely to find goal, must approximate.

# Optimal Decisions in Games

Game search problem:


- Initial state –
  - board position, player to move.
- Successor function –
  - returns list of {move, state} pairs, each indicating a legal move and the resulting state.
- Terminal test –
  - determines when the game is over, i.e., reach *terminal state*.
- Utility function (objective function) –
  - numeric value for terminal states  $\{+1, 0, -1\}$



# Optimal Decisions in Games

*Initial state* + the *legal moves* for each side define the *game tree*

# Game tree (2-player, deterministic, turns)



# Optimal Strategies

- Optimal solution is a sequence of moves leading to a goal state.
- Unfortunately, *Max* and *Min*, have equal and opposite objectives.
- Max must find a *contingent strategy*, which specifies Max's move in the initial state.
- Max moves in *successive states* resulting from every possible response by Min.
- *What approach can we use?*

# Minimax - Game Theory

- In the theory of *simultaneous games*, a ***minimax*** strategy is a *mixed strategy* which is part of the solution to a zero-sum game.
- In zero-sum games, the minimax solution is the same as the ***Nash equilibrium***:
  - A *solution concept* of a game involving two or more players.
  - Each player is assumed to know the equilibrium strategies of the other players.
  - No player has anything to gain by changing only his own strategy unilaterally.

# Minimax Theorem

- For every two-person, zero-sum game with finitely many strategies, there exists a value  $V$  and a mixed strategy for each player, such that:
  - Given player 2's strategy, the best payoff possible for player 1 is  $V$ , and
  - Given player 1's strategy, the best payoff possible for player 2 is  $-V$ .
- Equivalently, Player 1's strategy guarantees him a payoff of  $V$  regardless of Player 2's strategy, and similarly Player 2 can guarantee himself a payoff of  $-V$ .

# Minimax Theorem


- The name minimax arises because each player minimizes the maximum payoff possible for the other—since the game is zero-sum, he also maximizes his own minimum payoff.
- Theorem established by *John von Neumann*, who is quoted as saying:
  - "As far as I can see, there could be no theory of games ... without that theorem ... I thought there was nothing worth publishing until the *Minimax Theorem* was proved".

# Optimal Strategies (cont.)

- Optimal strategy leads to outcomes at *least* as good as any other strategy when one is playing an *infallible opponent*.
- Will do even better with fallible opponent.
- First, consider *optimal strategy*, even though its computationally problematic for all but the simplest games!
- Given a game tree, the optimal strategy can be determined by examining the *minimax* value of each node.

# Minimax

- Perfect play for deterministic games
- Idea: choose move with highest *minimax value*  
= best achievable payoff against best play
- E.g., 2-ply game:





# Minimax algorithm

**function** MINIMAX-DECISION(*state*) *returns an action*

$v \leftarrow \text{MAX-VALUE}(\textit{state})$

**return** the *action* in SUCCESSORS(*state*) with value *v*

---

**function** MAX-VALUE(*state*) *returns a utility value*

**if** TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

**for** *a, s* in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s))$

**return** *v*

---

**function** MIN-VALUE(*state*) *returns a utility value*

**if** TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow \infty$

**for** *a, s* in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s))$

**return** *v*

# Properties of minimax

- Complete? Yes! (if tree is finite)
- Optimal? Yes! (against an *optimal* opponent)
- Time complexity?  $O(b^m)$  – ouch!
- Space complexity?  $O(bm)$  (depth-first exploration)
- For chess,  $b \approx 35$ ,  $m \approx 100$  for "reasonable" games  
→ exact solution completely infeasible:  $35^{100}$
- *Note: minimax can easily be extended to  $> 2$  players.*

# Alpha-beta Pruning

**Huge** number of games states -  $O(b^m)$

- Can't eliminate exponent, but we can cut it in ~half.
- Trick:
  - Possible to compute the correct *minimax* decision without looking at *every* node in the game tree.
  - Using idea of pruning (from informed search) to eliminate (potentially) large parts of tree.
- **Alpha-beta** pruning
  - Apply to *minimax* game tree
  - Prune away branches that cannot possibly influence the final decision.

# Alpha-beta Pruning (cont.)

Consider simplification of the formula for minimax-value:

Minimax-value(root)


= max (min(3,12,8), min(2,x,y),min(14,5,2)) // leaf nodes

= max (3, min (2,x,y),2)


= max (3,z,2) where  $z \leq 2$

= 3


# $\alpha$ - $\beta$ pruning example




# $\alpha$ - $\beta$ pruning example



# $\alpha$ - $\beta$ pruning example




# $\alpha$ - $\beta$ pruning example





# $\alpha$ - $\beta$ pruning example



# Properties of $\alpha$ - $\beta$

- Pruning **does not** affect the final result
- Good move ordering improves effectiveness of pruning
- With "perfect ordering," time complexity =  $O(b^{m/2})$
- Example of the value of reasoning about which computations are relevant (*meta-reasoning*)

# $\alpha$ - $\beta$ ?

- $\alpha$  is the value of the best (i.e., highest-value) choice found so far at any choice point along the path for *max*
- If  $v$  is worse than  $\alpha$ , *max* will avoid it  
→ prune that branch
- Define  $\beta$  similarly for *min*

MAX

MIN


..

..

..

MAX

MIN



# The $\alpha$ - $\beta$ algorithm

**function** ALPHA-BETA-SEARCH(*state*) *returns an action*

**inputs:** *state*, current state in game

$v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$

**return** the *action* in SUCCESSORS(*state*) with value  $v$

---

**function** MAX-VALUE(*state*,  $\alpha$ ,  $\beta$ ) *returns a utility value*

**inputs:** *state*, current state in game

$\alpha$ , the value of the best alternative for MAX along the path to *state*

$\beta$ , the value of the best alternative for MIN along the path to *state*

**if** TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

**for**  $a, s$  in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s, \alpha, \beta))$

**if**  $v \geq \beta$  **then return**  $v$

$\alpha \leftarrow \text{MAX}(\alpha, v)$

**return**  $v$

# The $\alpha$ - $\beta$ algorithm

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value
  inputs: state, current state in game
            $\alpha$ , the value of the best alternative for MAX along the path to state
            $\beta$ , the value of the best alternative for MIN along the path to state

  if TERMINAL-TEST(state) then return UTILITY(state)
   $v \leftarrow +\infty$ 
  for  $a, s$  in SUCCESSORS(state) do
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s, \alpha, \beta))$ 
    if  $v \leq \alpha$  then return  $v$ 
     $\beta \leftarrow \text{MIN}(\beta, v)$ 
  return  $v$ 
```

# Resource limits

Suppose we have 100 seconds, explore  $10^4$  nodes/sec  
→  $10^6$  nodes per move

Standard approach:

- **cutoff test:**  
e.g., depth limit (perhaps add *quiescence search*)
- **evaluation function (heuristic)**  
= estimated desirability of position

# Resource limits

## Quiescence search –

- Select values unlikely to exhibit wild swings in value.
- E.g. chess - in positions where favorable captures can be made are *not* quiescent for an evaluation function that just counts material.
- Nonquiescent positions can be expanded further until quiescent positions are reached.

## Horizon affect –

- Arises when the program is facing a move by the opponent that causes serious damage and is ultimately unavoidable.
- Hard to avoid.
- E.g., pawn making it across the board & being replaced by queen.
- Can use *singular exclusions* – move that is clearly better than any other move at the current position. Can go beyond normal depth limit since its branching factor = 1. Very effective in chess.

# Evaluation functions

- For chess, typically **linear** weighted sum of **features**

$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

- e.g.,  $w_1 = 9$  with  
 $f_1(s) = (\text{number of white queens}) - (\text{number of black queens})$ , etc.
- Can use Machine Learning to optimize your *weights*.
- *More sophisticated – learn moves!*



# Cutting off search

*MinimaxCutoff* is identical to *MinimaxValue* except

1. *Terminal* is replaced by *Cutoff*
2. *Utility* is replaced by *Eval*

Does it work in practice?

$$b^m = 10^6, b=35 \rightarrow m=4$$

4-ply lookahead is a hopeless chess player!

- 4-ply  $\approx$  human novice
- 8-ply  $\approx$  typical PC, human master
- 12-ply  $\approx$  Deep Blue, Kasparov

# Deterministic games in practice

- Checkers: Chinook ended 40-year-reign of human world champion Marion Tinsley in 1994. Used a pre-computed endgame database defining perfect play for all positions involving 8 or fewer pieces on the board, a total of *444 billion* positions.
- Chess: Deep Blue defeated human world champion Garry Kasparov in a six-game match in 1997. Deep Blue searches *200 million* positions per second, uses very sophisticated evaluation, and undisclosed methods for extending some lines of search up to 40 ply.
- Othello: human champions refuse to compete against computers, who are too good.
- In 2008, thanks to an efficient message-passing parallelization, **MoGo** won one game (out of three) against Catalin Taranu, 5th dan pro, in 9x9 with standard time settings (30 minutes per side).



# Summary

- Reviewed games to understand what optimal play means and to understand how to play well in practice.
- Game can be defined by the *initial state* (how the board is set up), legal *actions* in each state, a *terminal test* (game over), and a *utility function* that applies to terminal states.
- For two-player, zero-sum games with perfect information, the *minimax* algorithm can select optimal moves using a depth-first enumeration of the game tree.
- *Alpha-beta* search computes the same optimal move as minimax, but achieves much greater efficiency by eliminating subtrees that are provably irrelevant.
- Usually not feasible to consider the whole game tree, need to *cut the search* off at some point and apply *eval* function that gives a *guesstimate* of the utility of a state.

# Summary

- Games of chance can be handled by an extension to minimax algorithm that values a *chance* node by taking the *average utility* of all its children nodes *weighted by the probability* of each child.
- Optimal play in games of imperfect information (bridge) requires reasoning about the current and future belief states of each player. Can average value of an action over each possible action.
- Programs can match or beat the best human players in many games.
- Games are fun to work on!
  - Illustrate several important points about AI
  - perfection is unattainable → must approximate
  - good idea to think about what to think about