# Logical Agents

CS4881 Artificial Intelligence

Jay Urbain, Ph.D.

Logical Agents

Outline

- Knowledge-based agents

- Wumpus world

- Logic in general - models and entailment

- Propositional (Boolean) logic

- Equivalence, validity, satisfiability

- Inference rules and theorem proving
  - forward chaining
  - backward chaining
  - resolution

Knowledge bases

- Knowledge base = set of sentences in a formal language

- Declarative approach to building an agent (or other system):
  - **Tell** it what it needs to know

- Then it can **Ask** itself what to do - answers should follow from the KB

- Agents can be viewed at the knowledge level
  i.e., what they know, regardless of how implemented

- Or at the implementation level
  - i.e., data structures in KB and algorithms that manipulate them

A simple knowledge-based agent

```
function KB-AGENT( percept) returns an action
    static: KB, a knowledge base
            t, a counter, initially 0, indicating time

    TELL(KB, MAKE-PERCEPT-SENTENCE( percept, t))
    action ← ASK(KB, MAKE-ACTION-QUERY(t))
    TELL(KB, MAKE-ACTION-SENTENCE(action, t))
    t ← t + 1
    return action
```

- The KB-agent must be able to:
  - Represent states, actions, etc.
  - Incorporate new percepts
  - Update internal representations of the world
  - Deduce hidden properties of the world
  - Deduce appropriate actions

Wumpus World PEAS description (Michael Genesereth)

- Performance measure
  - gold +1000, death -1000
  - -1 per step, -10 for using the arrow

- Environment
  - Squares adjacent to Wumpus are smelly
  - Squares adjacent to pits are breezy
  - Glitter *iff* gold is in the same square
  - Shooting kills Wumpus if you are facing it
  - Shooting uses up the only arrow
  - Grabbing picks up gold if in same square
  - Releasing drops the gold in same square

- Sensors: Stench, Breeze, Glitter, Bump, Scream
- Actuators: Left turn, Right turn, Forward, Grab, Release, Shoot
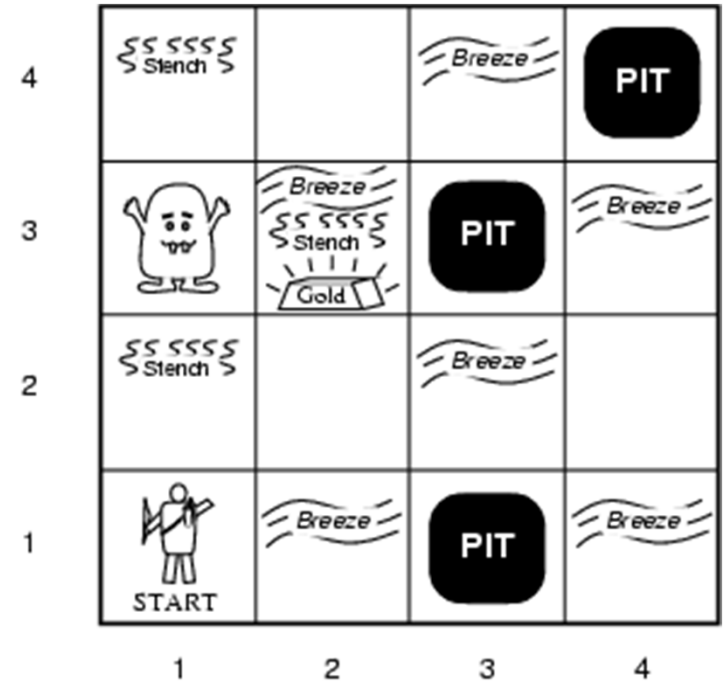
Wumpus world characterization

- <u>Fully Observable</u> No – only local perception

- <u>Deterministic</u> Yes – outcomes exactly specified

- <u>Episodic</u> No – sequential at the level of actions

- <u>Static</u> Yes – Wumpus and Pits do not move

- <u>Discrete</u> Yes

- <u>Single-agent?</u> Yes – Wumpus is essentially a natural feature

Exploring a Wumpus world

# Exploring a Wumpus world

# Exploring a Wumpus world

# Exploring a Wumpus world

Exploring a Wumpus world

Exploring a Wumpus world

Exploring a Wumpus world

Exploring a Wumpus world

Logic in general

- Logics are *formal languages* for representing information such that conclusions can be drawn

- Syntax defines the sentence structure in the language

- Semantics define the "meaning" of sentences;
  - i.e., define truth of a sentence in a world

- E.g., the language of arithmetic
  - x+2 ≥ y is a sentence; x2+y > {} is not a sentence
  - x+2 ≥ y is true iff the number x+2 is no less than the number y
  - x+2 ≥ y is true in a world where x = 7, y = 1
  - x+2 ≥ y is false in a world where x = 0, y = 6

Entailment

- Entailment means that one thing follows from another:

$$KB \models \alpha$$

- Knowledge base *KB* entails sentence α if and only if α is true in all worlds where *KB* is true

    – E.g., the KB containing "the Packers won" and "the Bears won" entails "Either the Packers won or the Bears won"
    – E.g., x+y = 4 entails 4 = x+y
    – E.g., x=2 and y=2 entails x+y = 4
    – Entailment is a relationship between sentences (i.e., syntax) that is based on semantics

Models

- Logicians typically think in terms of models, which are formally structured worlds with respect to which truth can be evaluated

- We say *m* is a model of a sentence α if α is true in *m*

- *M(α)* is the set of all models of α

- Then KB ⊨ α iff *M(KB)* ⊆ *M(α)*
  - E.g. *KB* = Packers won and Bears won α = Packers won

Entailment in the Wumpus world

Situation after detecting nothing in [1,1], moving right, breeze in [2,1]

Consider possible models for *KB* assuming only pits

3 Boolean choices $\Rightarrow$ 8 possible models

Wumpus models

Wumpus models

- *KB* = wumpus-world rules + observation

Wumpus models

- *KB* = wumpus-world rules + observations
- α_1 = "[1,2] is safe", $KB \models \alpha_1$, proved by model checking

Wumpus models

- *KB* = wumpus-world rules + observations
- α_2 = "[2,2] is safe", *KB* $\not\models \alpha_2$

Inference

- *KB* $\vdash_i$ α = sentence α can be derived from *KB* by procedure *i*
- Soundness: *i* is sound if whenever *KB* $\vdash_i$ α, it is also true that *KB* $\models$ α (only derives entailed sentences)
- Completeness: *i* is complete if whenever *KB* $\models$ α, it is also true that *KB* $\vdash_i$ α (can derive any entailed sentence)

- *Preview*: we will define a logic (first-order logic) which is expressive enough to say almost anything of interest, and for which there exists a sound and complete inference procedure.
- That is, the procedure will answer any question whose answer follows from what is known by the *KB*.

Propositional logic: Syntax

- Propositional (boolean) logic is the simplest logic – illustrates basic ideas

- The proposition symbols P_1, P_2, etc. are sentences

  - If S is a sentence, $\neg S$ is a sentence (negation)
  - If S_1 and S_2 are sentences, $S_1 \wedge S_2$ is a sentence (conjunction)
  - If S_1 and S_2 are sentences, $S_1 \vee S_2$ is a sentence (disjunction)
  - If S_1 and S_2 are sentences, $S_1 \Rightarrow S_2$ is a sentence (implication)
  - If S_1 and S_2 are sentences, $S_1 \Leftrightarrow S_2$ is a sentence (biconditional)

Propositional logic: Semantics

Each model specifies true/false for each proposition symbol

    E.g. $P_{1,2}$ $P_{2,2}$ $P_{3,1}$
         false   true   false

With these symbols, 8 possible models, can be enumerated automatically.

Rules for evaluating truth with respect to a model *m*:

        $\neg S$ is true iff   S is false
        $S_1 \wedge S_2$ is true iff S_1 is true and S_2 is true
        $S_1 \vee S_2$ is true iff S_1 is true or S_2 is true
        $S_1 \Rightarrow S_2$ is true iff S_1 is false or S_2 is true
         i.e.,   is false iff S_1 is true and S_2 is false
        $S_1 \Leftrightarrow S_2$ is true iff $S_1 \Rightarrow S_2$ is true and $S_2 \Rightarrow S_1$ is true

Simple recursive process evaluates an arbitrary sentence, e.g.,

$$\neg P_{1,2} \wedge (P_{2,2} \vee P_{3,1}) = \text{\textit{true}} \wedge (\text{\textit{true}} \vee \text{\textit{false}}) = \text{\textit{true}} \wedge \text{\textit{true}} = \text{\textit{true}}$$

Truth tables for connectives

P	Q	$\neg P$	$P \wedge Q$	$P \vee Q$	$P \Rightarrow Q$	$P \Leftrightarrow Q$
false	false	true	false	false	true	true
false	true	true	false	true	true	false
true	false	false	false	true	false	false
true	true	false	true	true	true	true

Wumpus world sentences

Let $P_{i,j}$ be true if there is a pit in [i, j].
Let $B_{i,j}$ be true if there is a breeze in [i, j].

$\quad \neg P_{1,1}$

$\quad \neg B_{1,1}$

$\quad B_{2,1}$

- "Pits cause breezes in adjacent squares"

$\quad B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})$

$\quad B_{2,1} \Leftrightarrow (P_{1,1} \vee P_{2,2} \vee P_{3,1})$

Truth tables for inference

$B_{1,1}$	$B_{2,1}$	$P_{1,1}$	$P_{1,2}$	$P_{2,1}$	$P_{2,2}$	$P_{3,1}$	KB	α_1
false	false	false	false	false	false	false	false	true
false	false	false	false	false	false	true	false	true
⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮
false	true	false	false	false	false	false	false	true
false	true	false	false	false	false	true	_true_	_true_
false	true	false	false	false	true	false	_true_	_true_
false	true	false	false	false	true	true	_true_	_true_
false	true	false	false	true	false	false	false	true
⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮
true	true	true	true	true	true	true	false	_false_

Inference by enumeration

- Depth-first enumeration of all models is sound and complete

```
function TT-ENTAILS?(KB, α) returns true or false
    symbols ← a list of the proposition symbols in KB and α
    return TT-CHECK-ALL(KB, α, symbols, [ ])

function TT-CHECK-ALL(KB, α, symbols, model) returns true or false
    if EMPTY?(symbols) then
        if PL-TRUE?(KB, model) then return PL-TRUE?(α, model)
        else return true
    else do
        P ← FIRST(symbols); rest ← REST(symbols)
        return TT-CHECK-ALL(KB, α, rest, EXTEND(P, true, model)) and
               TT-CHECK-ALL(KB, α, rest, EXTEND(P, false, model))
```

- For n symbols, time complexity is $O(2^n)$, space complexity is $O(n)$

Logical equivalence

- Two sentences are logically equivalent *iff* they are true in same models: $\alpha \equiv \beta$ iff $\alpha \models \beta$ and $\beta \models \alpha$

$$(\alpha \wedge \beta) \equiv (\beta \wedge \alpha) \quad \text{commutativity of } \wedge$$
$$(\alpha \vee \beta) \equiv (\beta \vee \alpha) \quad \text{commutativity of } \vee$$
$$((\alpha \wedge \beta) \wedge \gamma) \equiv (\alpha \wedge (\beta \wedge \gamma)) \quad \text{associativity of } \wedge$$
$$((\alpha \vee \beta) \vee \gamma) \equiv (\alpha \vee (\beta \vee \gamma)) \quad \text{associativity of } \vee$$
$$\neg(\neg\alpha) \equiv \alpha \quad \text{double-negation elimination}$$
$$(\alpha \Rightarrow \beta) \equiv (\neg\beta \Rightarrow \neg\alpha) \quad \text{contraposition}$$
$$(\alpha \Rightarrow \beta) \equiv (\neg\alpha \vee \beta) \quad \text{implication elimination}$$
$$(\alpha \Leftrightarrow \beta) \equiv ((\alpha \Rightarrow \beta) \wedge (\beta \Rightarrow \alpha)) \quad \text{biconditional elimination}$$
$$\neg(\alpha \wedge \beta) \equiv (\neg\alpha \vee \neg\beta) \quad \text{de Morgan}$$
$$\neg(\alpha \vee \beta) \equiv (\neg\alpha \wedge \neg\beta) \quad \text{de Morgan}$$
$$(\alpha \wedge (\beta \vee \gamma)) \equiv ((\alpha \wedge \beta) \vee (\alpha \wedge \gamma)) \quad \text{distributivity of } \wedge \text{ over } \vee$$
$$(\alpha \vee (\beta \wedge \gamma)) \equiv ((\alpha \vee \beta) \wedge (\alpha \vee \gamma)) \quad \text{distributivity of } \vee \text{ over } \wedge$$

Validity and satisfiability

A sentence is valid if it is true in all models,
e.g., *True*,   A $\vee \neg$A, A $\Rightarrow$ A, (A $\wedge$ (A $\Rightarrow$ B)) $\Rightarrow$ B

Validity is connected to inference via the Deduction Theorem:
KB $\models$ α if and only if (*KB* $\Rightarrow$ α) is valid

A sentence is satisfiable if it is true in some model
e.g., A$\vee$ B,   C

A sentence is unsatisfiable if it is true in no models
e.g., A$\wedge \neg$A

Satisfiability is connected to inference via the following:
KB $\models$ α if and only if (*KB* $\wedge \neg$α) is unsatisfiable

Proof methods

- Proof methods divide into (roughly) two kinds:

  - Application of inference rules
    - Legitimate (sound) generation of new sentences from old
    - Proof = a sequence of inference rule applications
    - Can use inference rules as operators in a standard search algorithm
    - Typically require transformation of sentences into a normal form

  - Model checking
    - truth table enumeration (always exponential in n)
    - improved backtracking, e.g., Davis-Putnam-Logemann-Loveland (DPLL)
    - heuristic search in model space (sound but incomplete)
              e.g., min-conflicts-like algorithms

Resolution

Conjunctive Normal Form (CNF)
   conjunction of disjunctions of literals, clauses

   E.g., $(A \lor \neg B) \land (B \lor \neg C \lor \neg D)$

- Resolution inference rule (for CNF):

$$\frac{\ell_i \lor \ldots \lor \ell_k, \qquad\qquad m_1 \lor \ldots \lor m_n}{\ell_i \lor \ldots \lor \ell_{i-1} \lor \ell_{i+1} \lor \ldots \lor \ell_k \lor m_1 \lor \ldots \lor m_{j-1} \lor m_{j+1} \lor \ldots \lor m_n}$$

where ℓ_i and m_j are complementary literals.
E.g., $\dfrac{P_{1,3} \lor P_{2,2}, \qquad \neg P_{2,2}}{P_{1,3}}$

- Resolution is sound and complete
  for propositional logic

Resolution

Soundness of resolution inference rule:

$$\neg(l_i \vee \ldots \vee l_{i-1} \vee l_{i+1} \vee \ldots \vee l_k) \Rightarrow l_i$$

$$\neg m_j \Rightarrow (m_1 \vee \ldots \vee m_{j-1} \vee m_{j+1} \vee \ldots \vee m_n)$$

$$\overline{\neg(l_i \vee \ldots \vee l_{i-1} \vee l_{i+1} \vee \ldots \vee l_k) \Rightarrow (m_1 \vee \ldots \vee m_{j-1} \vee m_{j+1} \vee \ldots \vee m_n)}$$

Conversion to CNF

$B_{1,1} \Leftrightarrow (P_{1,2} \lor P_{2,1})$

Conversion to CNF

$B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})$

1. Eliminate $\Leftrightarrow$, replacing $\alpha \Leftrightarrow \beta$ with $(\alpha \Rightarrow \beta) \wedge (\beta \Rightarrow \alpha)$.
   $(B_{1,1} \Rightarrow (P_{1,2} \vee P_{2,1})) \wedge ((P_{1,2} \vee P_{2,1}) \Rightarrow B_{1,1})$

2. Eliminate $\Rightarrow$, replacing $\alpha \Rightarrow \beta$ with $\neg\alpha \vee \beta$.
   $(\neg B_{1,1} \vee P_{1,2} \vee P_{2,1}) \wedge (\neg(P_{1,2} \vee P_{2,1}) \vee B_{1,1})$

3. Move $\neg$ inwards using de Morgan's rules and double-negation:
   $(\neg B_{1,1} \vee P_{1,2} \vee P_{2,1}) \wedge ((\neg P_{1,2} \vee \neg P_{2,1}) \vee B_{1,1})$

4. Apply distributivity law ($\wedge$ over $\vee$) and flatten:
   $(\neg B_{1,1} \vee P_{1,2} \vee P_{2,1}) \wedge (\neg P_{1,2} \vee B_{1,1}) \wedge (\neg P_{2,1} \vee B_{1,1})$

Resolution algorithm

- Proof by contradiction, i.e., show $KB \wedge \neg\alpha$ unsatisfiable

```
function PL-RESOLUTION(KB, α) returns true or false

    clauses ← the set of clauses in the CNF representation of KB ∧ ¬α
    new ← { }
    loop do
        for each Cᵢ, Cⱼ in clauses do
            resolvents ← PL-RESOLVE(Cᵢ, Cⱼ)
            if resolvents contains the empty clause then return true
            new ← new ∪ resolvents
        if new ⊆ clauses then return false
        clauses ← clauses ∪ new
```

Resolution example

- $KB = (B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})) \wedge \neg B_{1,1}$ $\alpha = \neg P_{1,2}$
- $(\neg B_{1,1} \vee P_{1,2} \vee P_{2,1}) \wedge (\neg P_{1,2} \vee B_{1,1}) \wedge (\neg P_{2,1} \vee B_{1,1}) \wedge \neg B_{1,1}$
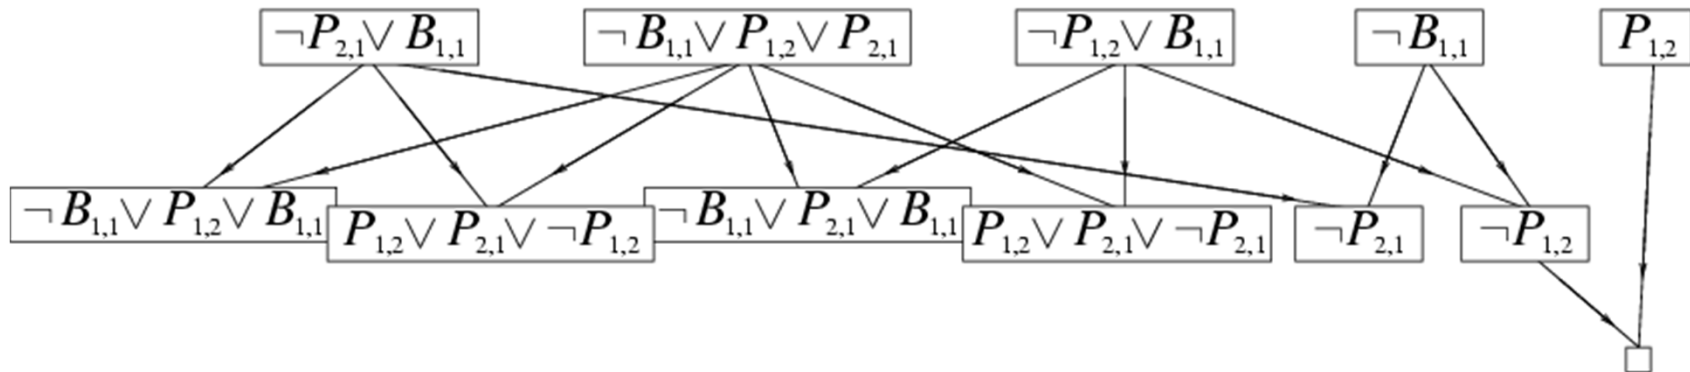
Forward and backward chaining

- Horn Form (restricted)

   KB = conjunction of Horn clauses

   – Horn clause =
      - proposition symbol; or
      - (conjunction of symbols) $\Rightarrow$ symbol
   – E.g., $C \wedge (B \Rightarrow A) \wedge (C \wedge D \Rightarrow B)$

- Modus Ponens (for Horn Form): complete for Horn KBs

$$\frac{\alpha_1 \wedge \ldots \wedge \alpha_n \Rightarrow \beta \qquad \alpha_1, \ldots, \alpha_n,}{\beta}$$

- Can be used with forward chaining or backward chaining.
- These algorithms are very natural and run in linear time

Forward chaining

- Idea: fire any rule whose premises are satisfied in the *KB*,
  - add its conclusion to the *KB*, until query is found

$$P \Rightarrow Q$$
$$L \wedge M \Rightarrow P$$
$$B \wedge L \Rightarrow M$$
$$A \wedge P \Rightarrow L$$
$$A \wedge B \Rightarrow L$$
$$A$$
$$B$$

Forward chaining algorithm

```
function PL-FC-ENTAILS?(KB, q) returns true or false
   local variables: count, a table, indexed by clause, initially the number of premises
                     inferred, a table, indexed by symbol, each entry initially false
                     agenda, a list of symbols, initially the symbols known to be true

   while agenda is not empty do
       p ← POP(agenda)
       unless inferred[p] do
           inferred[p] ← true
           for each Horn clause c in whose premise p appears do
               decrement count[c]
               if count[c] = 0 then do
                   if HEAD[c] = q then return true
                   PUSH(HEAD[c], agenda)
   return false
```

- Forward chaining is sound and complete for Horn KB

Forward chaining example

# Forward chaining example

# Forward chaining example

# Forward chaining example

Forward chaining example

Forward chaining example

Forward chaining example

Forward chaining example

Proof of completeness

- FC derives every atomic sentence that is entailed by *KB*

  1. FC reaches a fixed point where no new atomic sentences are derived

  2. Consider the final state as a model *m*, assigning true/false to symbols

  3. Every clause in the original *KB* is true in *m*

     $a_1 \wedge \ldots \wedge a_k \Rightarrow b$

  4. Hence *m* is a model of *KB*

  5. If $KB \models q$, *q* is true in every model of *KB*, including *m*

Backward chaining

Idea: work backwards from the query *q*:
 to prove *q* by BC,
 check if *q* is known already, or
  prove by BC all premises of some rule concluding *q*

Avoid loops: check if new sub-goal is already on the goal stack

Avoid repeated work: check if new sub-goal
  1. has already been proved true, or
  2. has already failed

Backward chaining example

Backward chaining example

Backward chaining example

Backward chaining example

Backward chaining example

Backward chaining example

Backward chaining example

Backward chaining example

Backward chaining example

Backward chaining example

Forward vs. backward chaining

- FC is data-driven, automatic, unconscious processing,
  - e.g., object recognition, routine decisions

- May do lots of work that is irrelevant to the goal

- BC is goal-driven, appropriate for problem-solving,
  - e.g., Where are my keys? How do I get into a PhD program?

- Complexity of BC can be much less than linear in size of KB

Efficient propositional inference

Two families of efficient algorithms for propositional inference:

Complete backtracking search algorithms

- DPLL algorithm (Davis, Putnam, Logemann, Loveland)
- Incomplete local search algorithms
  - `WalkSAT` algorithm

The DPLL algorithm

Determine if an input propositional logic sentence (in **CNF**) is satisfiable.

Improvements over truth table enumeration:
1. Early termination
   A clause is true if any literal is true.
   A sentence is false if any clause is false.

2. Pure symbol heuristic
   Pure symbol: always appears with the same "sign" in all clauses.
   e.g., In the three clauses $(A \lor \neg B)$, $(\neg B \lor \neg C)$, $(C \lor A)$, A and B are pure, C is impure.
   Make a pure symbol literal true.

3. Unit clause heuristic
   Unit clause: only one literal in the clause
   The only literal in a unit clause must be true.

The DPLL algorithm

function DPLL-SATISFIABLE?(s) **returns** *true* or *false*
    inputs: s, a sentence in propositional logic

    clauses ← the set of clauses in the CNF representation of s
    symbols ← a list of the proposition symbols in s
    return DPLL(*clauses*, *symbols*, [ ])

---

function DPLL(*clauses*, *symbols*, *model*) **returns** *true* or *false*

    if every clause in *clauses* is true in *model* **then return** *true*
    if some clause in *clauses* is false in *model* **then return** *false*
    P, *value* ← FIND-PURE-SYMBOL(*symbols*, *clauses*, *model*)
    if P is non-null **then return** DPLL(*clauses*, *symbols*–P, [$P = value | model$])
    P, *value* ← FIND-UNIT-CLAUSE(*clauses*, *model*)
    if P is non-null **then return** DPLL(*clauses*, *symbols*–P, [$P = value | model$])
    P ← FIRST(*symbols*);   *rest* ← REST(*symbols*)
    return DPLL(*clauses*, *rest*, [$P = true | model$]) **or**
           DPLL(*clauses*, *rest*, [$P = false | model$])

The `WalkSAT` algorithm

- Incomplete, local search algorithm
- Evaluation function: The min-conflict heuristic of minimizing the number of unsatisfied clauses
- Balance between greediness and randomness

The `WalkSAT` algorithm

function WALKSAT(*clauses*, p, *max-flips*) **returns** a satisfying model or *failure*
    inputs: *clauses*, a set of clauses in propositional logic
              p, the probability of choosing to do a "random walk" move
              max-flips, number of flips allowed before giving up

    model ← a random assignment of *true/false* to the symbols in *clauses*
    for $i = 1$ **to** *max-flips* **do**
        if *model* satisfies *clauses* **then return** *model*
        clause ← a randomly selected clause from *clauses* that is false in *model*
        with probability p flip the value in *model* of a randomly selected symbol
            from *clause*
      else flip whichever symbol in *clause* maximizes the number of satisfied clauses
    return *failure*

Hard satisfiability problems

- Consider random 3-CNF sentences. e.g.,

  ($\neg$D $\lor$ $\neg$B $\lor$ C) $\land$ (B $\lor$ $\neg$A $\lor$ $\neg$C) $\land$ ($\neg$C $\lor$ $\neg$B $\lor$ E) $\land$ (E $\lor$ $\neg$D $\lor$ B) $\land$ (B $\lor$ E $\lor$ $\neg$C)

  m = number of clauses

  n = number of symbols

  – Hard problems seem to cluster near m/n = 4.3 (critical point)

Hard satisfiability problems

Hard satisfiability problems

- Median runtime for 100 satisfiable random 3-CNF sentences, $n = 50$

Inference-based agents in the wumpus world

A wumpus-world agent using propositional logic:

$\neg P_{1,1}$
$\neg W_{1,1}$
$B_{x,y} \Leftrightarrow (P_{x,y+1} \lor P_{x,y-1} \lor P_{x+1,y} \lor P_{x-1,y})$
$S_{x,y} \Leftrightarrow (W_{x,y+1} \lor W_{x,y-1} \lor W_{x+1,y} \lor W_{x-1,y})$
$W_{1,1} \lor W_{1,2} \lor \ldots \lor W_{4,4}$
$\neg W_{1,1} \lor \neg W_{1,2}$
$\neg W_{1,1} \lor \neg W_{1,3}$
…

$\Rightarrow$ 64 distinct proposition symbols, 155 sentences

function PL-WUMPUS-AGENT(*percept*) **returns** an *action*
   inputs: *percept*, a list, $[stench, breeze, glitter]$
   static: KB, initially containing the "physics" of the wumpus world
        $x, y, orientation$, the agent's position (init. $[1,1]$) and orient. (init. *right*)
        visited, an array indicating which squares have been visited, initially *false*
        action, the agent's most recent action, initially null
        plan, an action sequence, initially empty

   update $x, y, orientation$, *visited* based on *action*
   if *stench* **then** TELL($KB, S_{x,y}$) **else** TELL($KB, \neg S_{x,y}$)
   if *breeze* **then** TELL($KB, B_{x,y}$) **else** TELL($KB, \neg B_{x,y}$)
   if *glitter* **then** *action* ← *grab*
   else if *plan* is nonempty **then** *action* ← POP(*plan*)
   else if for some fringe square $[i,j]$, ASK($KB, (\neg P_{i,j} \land \neg W_{i,j})$) is *true* **or**
        for some fringe square $[i,j]$, ASK($KB, (P_{i,j} \lor W_{i,j})$) is *false* **then do**
     $plan$ ← A*-GRAPH-SEARCH(ROUTE-PB($[x,y], orientation, [i,j], visited$))
     action ← POP(*plan*)
   else *action* ← a randomly chosen move
   return *action*

Expressiveness limitation of propositional logic

- KB contains "physics" sentences for every single square

- For every time t and every location $[x,y]$,

$$L_{x,y}^t \wedge FacingRight^t \wedge Forward^t \Rightarrow L_{x+1,y}^t$$

- Rapid proliferation of clauses

Summary

- Logical agents apply inference to a knowledge base to derive new information and make decisions
- Basic concepts of logic:
    - syntax: formal structure of sentences
    - semantics: truth of sentences with respect to models
    - entailment: necessary truth of one sentence given another
    - inference: deriving sentences from other sentences
    - soundness: derivations produce only entailed sentences
    - completeness: derivations can produce all entailed sentences
- Wumpus world requires the ability to represent partial and negated information, reason by cases, etc.
- Resolution is complete for propositional logic
  Forward, backward chaining are linear-time, complete for Horn clauses
- Propositional logic lacks expressive power